

← [Products](#) > [Deployment Orchestration](#) >
[Walmart Cloud Native Platform \(WCNP\)](#) > [Build and Deploy](#) > [Deployment](#) >
Kubectl

🔍 Search docs in Walmart Cloud Native Platf

Was this helpful? 👍 🗨

Kubectl

Last updated 6/28/2023

[Comments \(1\)](#)

[kubectl](#) is the command line tool that is used to interact with Kubernetes clusters, and therefore with your Managed Containers namespace.

After a local installation, the command line tool can be used from your workstation to perform mutating commands on non-production clusters.

Installation

You can download the binaries with curl from the external website as [documented](#). The remote repository is available internally as the proxy repository on Artifactory.

For example, the external URL

```
https://storage.googleapis.com/kubernetes-release/re
```

is available on Artifactory in the googleapis-storage repository at

Newer versions are available from the internal repository, even if they are not visible. Just update the version and request from the repository on the command-line.

```
curl -LO https://generic.ci.artifacts.walmart.com/artifa
```

A successful installation allows you to verify the version of kubectl:

```
$ kubectl version
Client Version: version.Info{Major:"1", Minor:"11", ... }
Server Version: version.Info{Major:"1", Minor:"9", ... }
```

The server version output only works after you have [configured kubectl](#).

Additionally, we recommend configuration of command completion for kubectl for bash or zsh users. This and other useful tips are documented on the [kubectl cheatsheet](#).

Installing Optional WCNP Plugins

To use these plugins, you should have already installed kubectl as shown above.

These plugins require kubectl 1.12.0 or higher, and a BASH execution environment.

Installing the plugins on a linux system is done using a single command:

```
curl "http://resources.k8s.walmart.com/get-kubectl-pl
```

Once installed, you may issue additional kubectl commands offered through the plugins, as described in the following sections.

Configuration

You can use [sledge](#) to configure kubectl access and connect it to your existing namespace.

```
sledge connect scus-lab-a1
```

Project-specific configuration can be stored in a k8s folder. It is then potentially available for [command line usage](#).

CLI Examples

Extensive documentation for command line usage of kubectl is available on the kubectl website, including a handy [cheat sheet](#).

As an application developer, there are a number of common issues for you to troubleshoot in your Kubernetes environment. Here are a number of samples that provide you with the ability to do a majority of the troubleshooting necessary on the Kubernetes side for most applications.

Using namespaces in kubectl commands

Your deployment should most likely be using a non-default [namespace](#).

To interact with your namespaced deployment, kubectl commands require you to provide the namespace as part of the command.

For example, if your kitt.yml contains namespace: strati-training, then you would run kubectl commands with a -n strati-training command line argument.

Check the status of your deployment

Use kubectl get deploy to view the status of all deployments, or include the name to check the status of a single deployment.

```
$ kubectl get deploy
NAME          DESIRED  CURRENT  UP-TO-DATE
frontend-deployment 3         3         0         5
```

Here there is a problem. Our deployment has the number of replicas that we want, but none of them are available for some reason. Let's take a closer look to see what might be occurring.

Check the detailed status of a deployment

Use `kubectl describe deploy {deployment name}` for a detailed view of a deployment and the overall status.

```
$ kubectl describe deploy frontend-deployment
Name:          frontend-deployment
Namespace:     default
CreationTimestamp: Tue, 09 Apr 2019 09:31:20 -0
Labels:        app=frontend
Annotations:    deployment.kubernetes.io/revision=1
                kubectl.kubernetes.io/last-applied-configuration={"apiVersion":"apps/v1","kind":"Deployment","metadata":{"name":"frontend-deployment","namespace":"default","uid":"1234567890","resourceVersion":"1234567890","creationTimestamp":"2019-04-09T09:31:20Z"},"spec":{"replicas":3,"selector":{"matchLabels":{"app":"frontend"},"matchExpressions":[{"key":"app","operator":"In","values":["nginx"]}]},"strategy":{"type":"RollingUpdate","rollingUpdate":{"maxUnavailable":"25%", "maxSurge":"25%"},"updateStrategyType":"RollingUpdate"},"template":{"metadata":{"labels":{"app":"nginx"},"annotations":{"kubectl.kubernetes.io/last-applied-configuration":"{\"apiVersion\":\"apps/v1\",\"kind\":\"Pod\",\"metadata\":{\"name\":\"frontend-deployment-5c689d88bb\",\"namespace\":\"default\",\"uid\":\"1234567890\",\"resourceVersion\":\"1234567890\",\"creationTimestamp\":\"2019-04-09T09:31:20Z\"},\"spec\":{\"containers\":[{\"name\":\"nginx\",\"image\":\"nginx:1.7.9\",\"ports\":[{\"containerPort\":80,\"protocol\":\"TCP\"}],\"hostPort\":0,\"protocol\":\"TCP\"}],\"environment\":[],\"mounts\":[],\"volumes\":[]}}]}\"}}"}
Selector:      app=nginx
Replicas:      3 desired | 3 updated | 3 total | 0 available
StrategyType:  RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=nginx
  Containers:
    nginx:
      Image:  nginx:1.7.9
      Port:   80/TCP
      Host Port: 0/TCP
      Environment: <none>
      Mounts:  <none>
      Volumes: <none>
  Conditions:
    Type             Status Reason
    ----             -
    Available        False MinimumReplicasUnavailable
    Progressing      False ProgressDeadlineExceeded
  OldReplicaSets: <none>
  NewReplicaSet:  frontend-deployment-5c689d88bb
Events:
  Type Reason      Age From          Message
  ---  -
  Normal ScalingReplicaSet 18m deployment-controller
```

This shows the status, the pod template used for the ReplicaSet, and the latest events. Pay attention to the `OldReplicaSet` and `NewReplicaSet` items. These contain the hash or the ReplicaSet currently being managed by the deployment. Using that, we can now view the individual pods to see whether we can find what might be wrong. Also, take note of the labels in the pod template. These will make finding the pods for this deployment a bit more easily in a large environment.

Check the status of a single pod in a deployment

First, we need to find the pods that are a part of this deployment. We can do that with `kubectl get pods` and the label mentioned in the deployment description above using the `-l` flag.

```
$ kubectl get pods -l app=nginx
NAME                                READY   STATUS    RES
frontend-deployment-5c689d88bb-6t4g6 0/1     Imaged
frontend-deployment-5c689d88bb-fn9k7 0/1     Imaged
frontend-deployment-5c689d88bb-vb84v 0/1     Imaged
```

This shows us the problem. Now let's go ahead and take a closer look at one of these pods to be certain we understand the issue.

We can view details of the pod, in the same way we did in the deployment by using the `describe` command and one of the pod names.

The pod name is using that ReplicaSet hash mentioned in the deployment description above. This is useful if you have large deployments that might be stuck in a rolling out status. Always make sure the hash matches the deployment Old or New ReplicaSet of interest.

```
$ kubectl describe pod frontend-deployment-5c689d88bb-6t
Name:          frontend-deployment-5c689d88bb-6t
Namespace:     default
Priority:       0
PriorityClassName: <none>
Node:          worke1547000001/10.74.248.40
Start Time:    Tue, 09 Apr 2019 09:31:21 -0500
Labels:        app=nginx
               pod-template-hash=5c689d88bb
Annotations:   <none>
Status:        Pending
IP:            10.24.21.231
Controlled By: ReplicaSet/frontend-deployment-5c689d88bb
Containers:
  nginx:
    Container ID:
    Image:        nginx:1.7.9
    Image ID:
    Port:         80/TCP
    Host Port:    0/TCP
    State:        Waiting
      Reason:     ImagePullBackOff
    Ready:        False
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from frontend-deployment-token-b6qsz (ro)
Conditions:
  Type          Status
  Initialized    True
  Ready          False
  ContainersReady False
```

[Get Started](#)[Products](#)[Blog](#)[Stack Overflow](#)[Feedback](#)[Help](#)[Console](#)[YJ](#)

Walmart Cloud Native Platform (WCNP)

[Deploying Applications](#)[sledge](#)[Monitoring and Alerting](#)[GSLB Routing](#)[sledge kitt](#)[Kittbot](#)[Cronjob Helm Chart](#)[Secrets](#)[Kubectl](#)[Flagger Configurations](#)[HashiCorp Vault Secrets](#)[KITT Deployment Guide](#)[Namespaces](#)[Basic Web Application Helm Chart](#)[Application Configuration](#)[Deployment](#)[Learn more about GenAI Development Tracks - \[View Overview.\]\(#\)](#)

```
type: Secret (a volume populated by secrets)
SecretName: default-token-b6qsz
Optional: false
QoS Class:   BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute
              node.kubernetes.io/unreachable:NoExecute
Events:
  Type    Reason      Age    From          Message
  ----    -
  Normal  Scheduled   30m    default-schedule
  Normal  Pulling     28m (x4 over 30m)  kubelet, wor
  Warning Failed      28m (x4 over 30m)  kubelet, wor
  Warning Failed      28m (x4 over 30m)  kubelet, wor
  Warning Failed      5m39s (x99 over 30m) kubelet, w
  Normal  BackOff     41s (x120 over 30m) kubelet, w
```

This gives us much more information. Specifically the last few events tell us what problem is causing the image pull failure. We can't connect to the docker registry directly. Because these systems live inside of the Walmart network, we have to use the local artifactory repos for accessing outside resources. For more details, go to the [KITT Docker documentation](#).

We can make the changes by editing the deployment and checking the status.

```
$ kubectl get deploy
NAME                DESIRED  CURRENT  UP-TO-DATE  AVAILABLE
frontend-deployment 3          3          3          3
```

TABLE OF CONTENTS

Installation

[Installing Optional WCNP Plugins](#)[Configuration](#)[CLI Examples](#)[Using namespaces in kubectl commands](#)[Check the status of your deployment](#)[Check the detailed status of a deployment](#)[Check the status of a single pod in a deployment](#)[Viewing the application logs inside of a pod](#)[Getting a command line in your container](#)[Perform a manual scaling request](#)[Perform a rolling restart](#)

Viewing the application logs inside of a pod

Now that our application is running in a deployment, we will need to view the actual application logs in the container. While these logs are most likely available in your Splunk or ELK instance, it can be valuable to look at the actual container logs, especially if you are having an issue with log collection.

To view the logs of a pod with only one container you can use `kubectl logs {podname}`:

```
$ kubectl logs frontend-deployment-5d7c89c76c-djg
80.91.33.133 - - [Day/Month/Year:08:05:50 +0000]
173.203.139.108 - - [Day/Month/Year:08:05:03 +00
80.91.33.133 - - [Day/Month/Year:08:05:35 +0000]
5.83.131.103 - - [Day/Month/Year:08:05:51 +0000]
80.91.33.133 - - [Day/Month/Year:08:05:59 +0000]
200.6.73.40 - - [Day/Month/Year:08:05:42 +0000] "C
80.91.33.133 - - [Day/Month/Year:08:05:48 +0000] "C
93.180.71.3 - - [Day/Month/Year:08:05:58 +0000] "C
```

If you have a pod with multiple containers you will need to use the `-c` flag to specify which containers logs you want to view.

Getting a command line in your container

In some cases, it may be necessary to get a command line inside of your actual container to troubleshoot an issue you are having. Provided your base image is Alpine, or any other distribution with a full shell, you can connect to a container and execute a command.

Using `kubectl exec -it {pod name} -- /bin/bash` you can execute a bash shell on a particular pod. You can also execute commands directly when you know their location, or they are part of the containers internal path.

```
$ kubectl exec -it frontend-deployment-5d7c89c76c-
nginx version: nginx/1.11.4
built by gcc 4.9.2 (Debian 4.9.2-10)
built with OpenSSL 1.0.1t  3 May 2016
TLS SNI support enabled
configure arguments: --prefix=/etc/nginx --sbin-path=
```

Perform a manual scaling request

_This action results in a modification to cluster objects outside of the normal KITT release process.

Depending on the nature of other changes in the kitt file, the changes made here may persist or get overridden from the kitt file.

For example, a manual update of `HorizontalPodAutoscaler` `minReplicas` or `maxReplicas` using the `kubectl` will ignore the `kitt.yml` `scaling.min` or `scaling.max` values. To retain any scaling configuration, you must go back and update the corresponding `kitt.yml` in your GitHub repo._

Use of this requires installing WCNP `kubectl` plugins as described [above](#).

While KITT-managed deployments are set up by default to use autoscaling, there may be times during larger-scale testing, or one-time predicted traffic spikes, where

manual scale up is desired.

This command performs a manual scale up or scale down of the provisioned HPA for the deployment.

Example Usage:

```
$ kubectl scale-hpa --context=scus-lab-a1 --hpa myapp scaling hpa/myapp...
horizontalpodautoscaler.autoscaling/myapp configured
hpa myapp successfully updated to min=2 and max=3.
$
```

More information about the command is available using the `kubectl help` option for the command.

Perform a rolling restart

Use of this requires installing WCNP kubectl plugins as [described above](#).

Deleting a pod serves as a restart for a particular app instance. But there may be times when a restart of all running instances is desired. If the application has many instances, it can be cumbersome to issue individual pod delete commands for every instance.

This command performs a *rolling restart* of all active pods for a deployment using a single convenient command. Rolling restarts help avoid app outages by keeping one or more active pods up. The number of restarting pods is determined by the metadata settings of the deployment itself. With KITT this is defaulted to a maximum of 25% of the available pods.

Example Usage:

```
$ kubectl rolling-restart --timeout 3m --context=scus-lab-a1
starting rolling restart of deployment/myapp...
deployment.extensions/myapp env updated
Waiting for deployment "myapp" rollout to finish: 1 old pods
Waiting for deployment "myapp" rollout to finish: 1 old pods
deployment "myapp" successfully rolled out
Finished rolling restart of deployment myapp.
$
```

The timeout option, `-t` or `--timeout`, is used purely for polling purposes, after the restart command has been submitted to the cluster. If the polling stops before the cluster operation is complete, the restarts will continue on the cluster, up until `progressDeadlineSeconds` is exceeded. The value for `progressDeadlineSeconds` is set by KITT to a default of 600 on the kubernetes deployment at the time of app installation.

More information about the command is available using the `kubectl help` option for the command.

Was this helpful?  

Comments(1)

YJ

LG **Lin Gong** May 18th 2022, 3:04 PM

Use homebrew to install is much easier:
brew install kubect!
(edited)

  26 